

Machine Learning 01


Meg Robinson

2/14/2022

Clustering Methods

First, create some data to test and learn with

```
tmp = c(rnorm(30,3), rnorm(30,-3))  
data = cbind(x=tmp, y= rev(tmp))  
plot(data)
```



Run `kmeans()`, specifying # of clusters.

```
k = kmeans(data,centers=2,nstart=20)  
k  
  
## K-means clustering with 2 clusters of sizes 30, 30  
##  
## Cluster means:  
##      x      y  
## 1 -2.871373  3.151390  
## 2  3.151390 -2.871373
```

```
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1
1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##
## Within cluster sum of squares by cluster:
## [1] 87.8321 87.8321
## (between_SS / total_SS = 86.1 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"
"tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

QUESTION: How many points are in each cluster?

```
k$size
```

```
## [1] 30 30
```

ANSWER: 30 points in each cluster

QUESTION: What component of your result object details cluster size?

```
k$cluster
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1
1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


QUESTION: What component of your result object details cluster center?

```
k$centers
```

```
##           x           y
## 1 -2.871373  3.151390
## 2  3.151390 -2.871373
```

QUESTION: Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(data, col=k$cluster)
points(k$centers, col='blue', pch=15, cex=1.5)
```



Hierarchical Clustering


```
h = hclust(dist(data))
h

##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

Use the plot method:

```
plot(h)
abline(h=7,col='red')
```

Cluster Dendrogram



dist(data)
hclust (*, "complete")


Find the membership vector using cutree()

```
cutree(h, h=7)

## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

And use this method to find number of k clusters

```
kclst = cutree(h, k=2)
plot(data, col=kclst)
```



We can see that `kmeans()` uses the data and number of centers, while `hclust()` uses the distance of the data.

PCA of UK food data

import data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named `x`? What R functions could you use to answer this questions?

Complete the following code to find out how many rows and columns are in `x`?

```
__(x)
```

```
dim(x)
```

```
## [1] 17 5
```

Preview the first 6 rows

```
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105    103      103        66
## 2 Carcass_meat      245    227      242       267
```

```
## 3    Other_meat      685    803      750      586
## 4          Fish      147    160      122       93
## 5 Fats_and_oils      193    235      184      209
## 6          Sugars      156    175      147      139
```

Row titles are incorrectly being stored as a column. Fix this.

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)

##           England Wales Scotland N.Ireland
## Cheese           105    103      103       66
## Carcass_meat      245    227      242      267
## Other_meat        685    803      750      586
## Fish             147    160      122       93
## Fats_and_oils     193    235      184      209
## Sugars            156    175      147      139

dim(x)

## [1] 17  4

x <- read.csv(url, row.names=1)
head(x)

##           England Wales Scotland N.Ireland
## Cheese           105    103      103       66
## Carcass_meat      245    227      242      267
## Other_meat        685    803      750      586
## Fish             147    160      122       93
## Fats_and_oils     193    235      184      209
## Sugars            156    175      147      139
```


Now there are correctly 17 rows and 4 columns, yay!

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second approach, `read.csv()` is way better than the first one in general because it loads the data without manipulation. The first approach, indexing `x[, -1]`, removes data & will continue to remove the last column every time it is run.

Spotting major differences and trends


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


We can remove the `beside=TRUE` argument to obtain the following:

```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```

The y axis for each row of plots is the country name in that row, whereas the x axis for each column of the plots is the country name in that column. Therefore, the diagonal plots of this graph indicates that the country is the same on the row and column so there is no way to plot a pairwise interaction.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland appears to be more unique as it contains the most values that deviate from the diagonal when compared to the other countries.

PCA to the rescue

Use the `prcomp()` PCA function

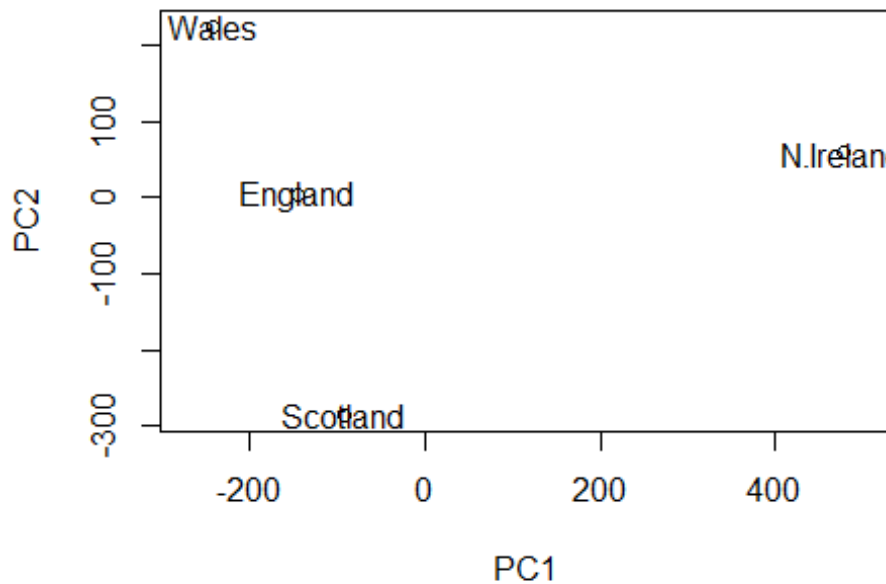
```
pca <- prcomp( t(x) )
summary(pca)

## Importance of components:
##               PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

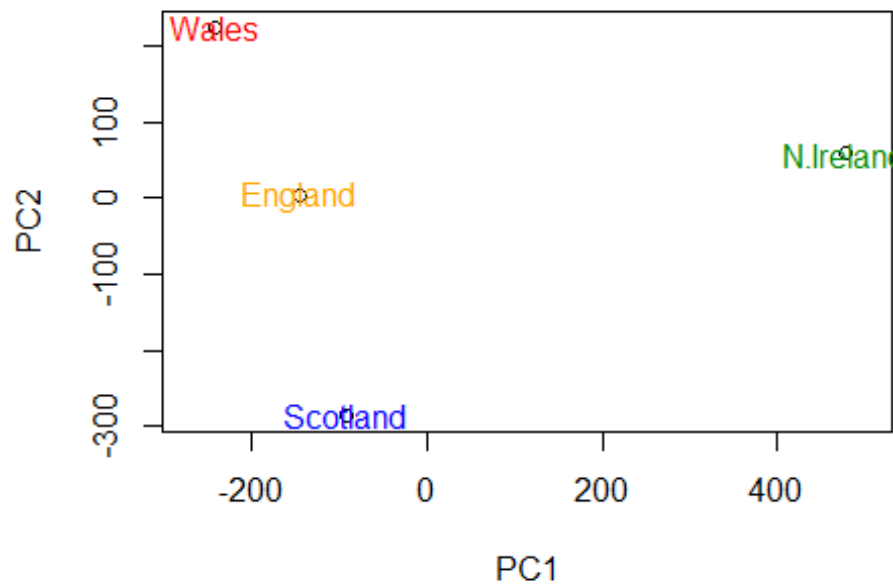
Plot PC1 vs PC2

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))  
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
mycols = c("orange", "red", "blue", "green4")  
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))  
text(pca$x[,1], pca$x[,2], colnames(x), col=mycols)
```




```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
## [1] 67 29 4 0

## or the second row here...
z <- summary(pca)
z$importance

##               PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00

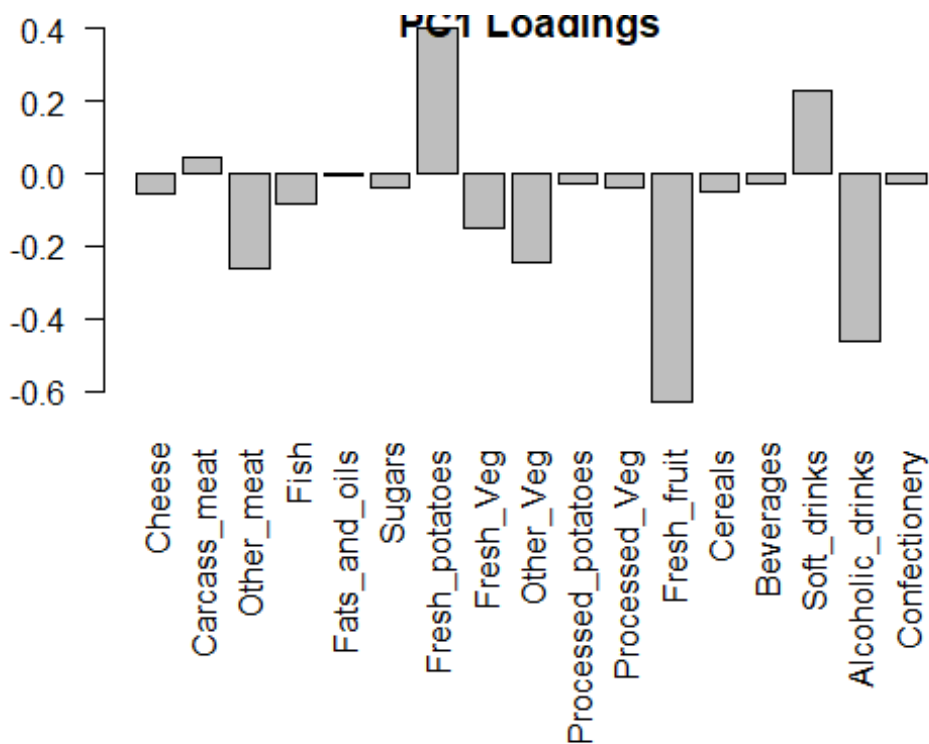
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



> Digging deeper


(variable loadings):

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2, main="PC1 Loadings" )
```



> Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?


```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2, main="PC2 Loadings" )
```



From looking at this graph, it looks like 'fresh potatoes' and 'soft drinks' feature prominently. The plot also tells us which groups contribute the most to the remaining 10% variance, after PC1.

Biplots

The inbuilt biplot() can be useful for small datasets
 biplot(pca)



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

```
nrow(rna.data)

## [1] 100

ncol(rna.data)


## [1] 10
```

There are 100 genes and 10 samples.

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")

text(pca$x[,1:2], labels = colnames(rna.data))
```




```
summary(pca)

## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
PC7
## Standard deviation    9.6237  1.5198  1.05787  1.05203  0.88062  0.82545
0.80111
## Proportion of Variance 0.9262  0.0231  0.01119  0.01107  0.00775  0.00681
0.00642
## Cumulative Proportion 0.9262  0.9493  0.96045  0.97152  0.97928  0.98609
0.99251
##              PC8      PC9      PC10
## Standard deviation    0.62065  0.60342  3.348e-15
## Proportion of Variance 0.00385  0.00364  0.000e+00
## Cumulative Proportion 0.99636  1.00000  1.000e+00

plot(pca, main="Quick scree plot")
```


Quick scree plot




Lets make our own scree plots:

```
## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per

## [1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0

barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```



Make it more attractive and useful:

A vector of colors for wt and ko samples

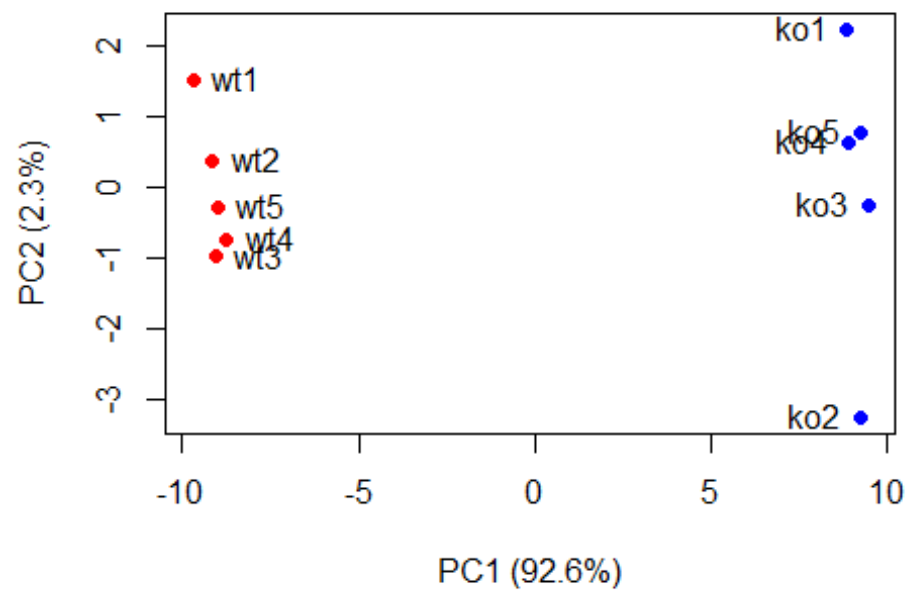
```
colvec <- colnames(rna.data)
```

```
colvec[grep("wt", colvec)] <- "red"
```

```
colvec[grep("ko", colvec)] <- "blue"
```

```
plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))
```

```
text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5),
rep(2,5)))
```




Using ggplot

```
library(ggplot2)


df <- as.data.frame(pca$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)


p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="BIMM143 example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



BIMM143 example data